

```

1 !pip install kaggle
2 !pip install seaborn
3
4 import os
5
6 # Set the path to the Kaggle API credentials
7
8 os.environ['KAGGLE_CONFIG_DIR'] = "/content/.kaggle"
9
10 !kaggle datasets download -d aryashah2k/breast-ultrasound-images-dataset # Download
11
12 !unzip breast-ultrasound-images-dataset.zip # unzip
13

```

```

➞ Requirement already satisfied: kaggle in /usr/local/lib/python3.10/dist-packages (1.5.12)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.10/dist-packages (1.16.0)
Requirement already satisfied: certifi>=2023.7.22 in /usr/local/lib/python3.10/dist-packages (2023.11.17)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.10/dist-packages (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (2.31.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (4.66.1)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.10/dist-packages (8.0.4)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (2.0.7)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (6.1.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (0.11)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.10/dist-packages (1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (3.6)
Requirement already satisfied: seaborn in /usr/local/lib/python3.10/dist-packages (0.12.2)
Requirement already satisfied: numpy!=1.24.0,>=1.20 in /usr/local/lib/python3.10/dist-packages (1.26.4)
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.10/dist-packages (2.2.2)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in /usr/local/lib/python3.10/dist-packages (3.8.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (1.1.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (4.52.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (24.0)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (10.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (2023.3)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (2023.3)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (1.16.0)
Dataset URL: https://www.kaggle.com/datasets/aryashah2k/breast-ultrasound-images-dataset
License(s): CC0-1.0
breast-ultrasound-images-dataset.zip: Skipping, found more recently modified
Archive: breast-ultrasound-images-dataset.zip
replace Dataset_BUSI_with_GT/benign/benign (1).png? [y]es, [n]o, [A]ll, [N]one
  inflating: Dataset_BUSI_with_GT/benign/benign (1).png
  inflating: Dataset_BUSI_with_GT/benign/benign (1)_mask.png
  inflating: Dataset_BUSI_with_GT/benign/benign (10).png
  inflating: Dataset_BUSI_with_GT/benign/benign (10)_mask.png
  inflating: Dataset_BUSI_with_GT/benign/benign (100).png
  inflating: Dataset_BUSI_with_GT/benign/benign (100)_mask.png

```

```

inflating: Dataset_BUSI_with_GT/benign/benign (100)_mask.png
inflating: Dataset_BUSI_with_GT/benign/benign (100)_mask_1.png
inflating: Dataset_BUSI_with_GT/benign/benign (101).png
inflating: Dataset_BUSI_with_GT/benign/benign (101)_mask.png
inflating: Dataset_BUSI_with_GT/benign/benign (102).png
inflating: Dataset_BUSI_with_GT/benign/benign (102)_mask.png
inflating: Dataset_BUSI_with_GT/benign/benign (103).png
inflating: Dataset_BUSI_with_GT/benign/benign (103)_mask.png
inflating: Dataset_BUSI_with_GT/benign/benign (104).png
inflating: Dataset_BUSI_with_GT/benign/benign (104)_mask.png
inflating: Dataset_BUSI_with_GT/benign/benign (105).png
inflating: Dataset_BUSI_with_GT/benign/benign (105)_mask.png
inflating: Dataset_BUSI_with_GT/benign/benign (106).png
inflating: Dataset_BUSI_with_GT/benign/benign (106)_mask.png
inflating: Dataset_BUSI_with_GT/benign/benign (107).png
inflating: Dataset_BUSI_with_GT/benign/benign (107)_mask.png
inflating: Dataset_BUSI_with_GT/benign/benign (108).png
inflating: Dataset_BUSI_with_GT/benign/benign (108)_mask.png
inflating: Dataset_BUSI_with_GT/benign/benign (109).png
inflating: Dataset_BUSI_with_GT/benign/benign (109)_mask.png

```

We import the dataset and see that there is quite a lot of files to organize. There are subfolders that correspond to categories (benign, normal, malignant). We have to organize the data and put it into a format that can be manipulated before we can clean or preprocess the data.

```

1 # Import the os library
2
3 import os
4 from sklearn.metrics import f1_score, precision_score, recall_score, accuracy_
5 import tensorflow as tf
6 import numpy as np
7
8
9 # Change image_dir to the unzipped dataset path
10
11 image_dir = '/content/Dataset_BUSI_with_GT'
12 benign = '/content/Dataset_BUSI_with_GT/benign'
13 malignant = '/content/Dataset_BUSI_with_GT/malignant'
14 normal = '/content/Dataset_BUSI_with_GT/normal'
15 images = [benign, malignant, normal]
16

```

```

1 # Get all files in the directory
2 image_files = []
3 for folder_path in images:
4     files = [f for f in os.listdir(folder_path) if f.endswith(('.png', '.jpg', '
5     image_files.extend(files)
6 print(image_files)

```

```


['benign (125)_mask.png', 'benign (239)_mask.png', 'benign (378)_mask.png', 'I

```

The array `image_files` stores the names of all the image files. Names look like 'benign (125)\_mask.png', which are added by a for loop checking for the appropriate file extension in the current `folder_path`. The list is then printed to check progress.

```
1 # Display all images
2 from tensorflow import keras
3 import matplotlib.pyplot as plt
4
5 for image_file in image_files[:5]:
6     # Get the correct folder name from the 'images' list
7     for folder_path in images:
8         if image_file in os.listdir(folder_path):
9             folder_name = os.path.basename(folder_path)
10            break
11
12    image_path = os.path.join(folder_path, image_file) # Correct the image path
13    image = keras.utils.load_img(image_path)
14    # Convert image to display - matplotlib
15    plt.imshow(image)
16    plt.title(image_file)
17    plt.axis('off') # Turn off axis labels
18    plt.show()
```

benign (125)\_mask.png




benign (239)\_mask.png





benign (378)\_mask.png



benign (419)\_mask.png





benign (346)\_mask\_1.png



Outer for loop: displays the first five images in the list (image\_files). Inner for loop and if statment: find correct folder path for the image in the image\_file list. The code then joins the image file name and the folder path with image\_path = os.path.join(folder\_path, image\_file). The object image is created and the image is loaded usint keras.utils.load\_img. The image is loaded but not yet displayed until plt.imshow(image) is called using matplotlib.pyplot.imshow. The title is set as the name of the image (image\_file)

```
1 image_dir = '/content/Dataset_BUSI_with_GT'
2 subfolders = ["benign", "malignant", "normal"]
```

```
1 image_path = []
2 image_label = []
3 image_data = []
4 for subfolder in subfolders:
5     subfolder_path = os.path.join(image_dir, subfolder)
6     # Get all image files in the subfolder
```

```

7 image_files = [f for f in os.listdir(subfolder_path) if f.endswith(('.png',
8 for i in image_files:
9     image_path.append(os.path.join(subfolder_path, image_file))
10    image_label.append(subfolder)
11    image_data.append(image_file)
12
13
14 len(image_label)
15 len(image_data)

```

1578

```

1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(image_data, image_label, t

```

The code begins to use the path to the image and the subfolders. So far, we are still preprocessing. The nested for loop iterates through the subfolders and image files to populate the lists: image\_path, image\_label, and image\_data. The lengths of these lists are checked against the lengths of the label and data lists. Before we even start preprocessing, we split the data. As always we use a random\_state seed to make our split random and our results more reproducible.

```

1 import numpy as np
2 X_train = np.array(X_train)
3 print(X_train.shape)
4 y_train = np.array(y_train)
5 print(y_train.shape)
6 X_test = np.array(X_test)
7 print(X_test.shape)
8 y_test = np.array(y_test)
9 print(y_test.shape)
10

```

```

(1262,)
(1262,)
(316,)
(316,)

```

```

1 image_files

```

```

['normal (92)_mask.png',
 'normal (94).png',
 'normal (16)_mask.png',
 'normal (74)_mask.png',
 'normal (73)_mask.png',
 'normal (129)_mask.png',
 'normal (51).png']

```

```

'normal (45)_mask.png',
'normal (19).png',
'normal (46).png',
'normal (120)_mask.png',
'normal (34).png',
'normal (18)_mask.png',
'normal (117).png',
'normal (60)_mask.png',
'normal (97).png',
'normal (34)_mask.png',
'normal (89)_mask.png',
'normal (58)_mask.png',
'normal (106).png',
'normal (33)_mask.png',
'normal (73).png',
'normal (90)_mask.png',
'normal (85)_mask.png',
'normal (113).png',
'normal (68)_mask.png',
'normal (3).png',
'normal (20).png',
'normal (10).png',
'normal (81)_mask.png',
'normal (30).png',
'normal (102)_mask.png',
'normal (17)_mask.png',
'normal (95)_mask.png',
'normal (3)_mask.png',
'normal (17).png',
'normal (7).png',
'normal (44).png',
'normal (87).png',
'normal (16).png',
'normal (119)_mask.png',
'normal (60).png',
'normal (37)_mask.png',
'normal (5)_mask.png',
'normal (86)_mask.png',
'normal (27).png',
'normal (96)_mask.png',
'normal (30)_mask.png',
'normal (113)_mask.png',
'normal (85).png',
'normal (93)_mask.png',
'normal (99).png',
'normal (61).png',
'normal (79).png',
'normal (109).png',
'normal (108)_mask.png',
'normal (98).png',
'normal (78).png',

```

```

1 import tensorflow as tf
2
3 def f1_matrix(x, t, y, y_pred):

```

```

3 def f1_metric(y_true, y_pred):
4     y_true = tf.cast(y_true, tf.float32)
5     y_pred = tf.cast(tf.round(y_pred), tf.float32)
6     tp = tf.reduce_sum(y_true * y_pred)
7     fp = tf.reduce_sum((1 - y_true) * y_pred)
8     fn = tf.reduce_sum(y_true * (1 - y_pred))
9
10    precision = tp / (tp + fp + tf.keras.backend.epsilon())
11    recall = tp / (tp + fn + tf.keras.backend.epsilon())
12    f1 = 2 * (precision * recall) / (precision + recall + tf.keras.backend.epsilon())
13    return f1

```

```

1 #create an object for the model
2 cnn = tf.keras.models.Sequential()

```

```

1 # initialize the model and first hidden layer
2 cnn.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, activation='relu', input_shape=(3, 32, 32)))

```

`/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv2d.py:100: Warning: 'input_shape' argument is deprecated, use 'input_shape' attribute of the layer instead.`

```

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```

1 # max pooling - encoding
2 cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))

```

```

1 #second hidden layer and max pooling
2 cnn.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, activation='relu'))
3 cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))
4
5 cnn.summary()

```

**Model: "sequential"**

Layer (type)	Output Shape	
conv2d (Conv2D)	(None, 62, 62, 32)	
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	
conv2d_1 (Conv2D)	(None, 29, 29, 32)	
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 32)	

**Total params:** 10,144 (39.62 KB)  
**Trainable params:** 10,144 (39.62 KB)  
**Non-trainable params:** 0 (0.00 B)

```

1 # flatten dimensionality reduction
2 cnn.add(tf.keras.layers.Flatten())

```



```
3 cnn.add(tf.keras.layers.Dropout(0.25))
```

```
1 # third hidden layer
2 cnn.add(tf.keras.layers.Dense(units=128, activation='relu'))
```

```
1 # fourth hidden layer
2 cnn.add(tf.keras.layers.Dense(units=3, activation='softmax'))
```

```
1 # compile with adam as the optimizer, and accuracy as the metric
2 cnn.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = [
```

In the code above we use the TensorFlow library to an object for the model, add layers to the CNN model.

Layers:

- Conv2d: convolutional layers for feature extraction
- MaxPool2D: Max pooling layers for downsampling
- Flatten: Flattens the output for the dense layers
- Dropout: drops some of the nodes randomly to decrease risk of overfitting
- Dense: Fully connects layers for classification

Then we use `cnn.summary()` to see the model layers we have made also called the model architecture. Finally we compile the model with 'adam' as the optimizer

We are using f1 metric that is calculated directly using Tensorflow operations. This avoids any trouble shooting with numpy arrays within the metric function. We add `tf.keras.backend.epsilon()` to avoid dividing by zeros

```
1 # Create an ImageDataGenerator for training data
2 from tensorflow.keras.preprocessing.image import ImageDataGenerator
3 from tensorflow.keras.callbacks import ReduceLROnPlateau
4 train_datagen = ImageDataGenerator(rescale=1./255, # Rescale pixel values
5                                   shear_range=0.2,
6                                   zoom_range=0.2,
7                                   horizontal_flip=True,
8                                   rotation_range = 20,
9                                   width_shift_range = 0.2,
10                                  height_shift_range = 0.2,
11                                  brightness_range = (0.8, 1.2)
12                                  )
```

ImageDataGenerator is something new that came up as we started to build neural networks. ImageDataGenerator is assigned to the object `train_datagen`.

ImageDataGenerator augments parameters such as:

- `rescale = 1./225`: Rescales pixel values with a range of [0,1]
- `shear_range = 0.2`: shearing is tilting the image, so the tilting direction is randomized at 20%.
- `zoom_range=0.2`: the zooming transformation is also randomized at 20%.
- `horizontal_flip=True`: randomizes the flipping of images in a horizontal direction.
- `rotation_range=20`: rotates 20 degrees randomly
- `width_shift_range=0.2`: 20% width shift randomly
- `height_shift_range=0.2`: 20% height shift randomly

The overall purpose of the ImageDataGenerator is to create variations of the images to prevent overfitting. These transformations also increase the size of the training set, which is sometimes very necessary for small datasets.

This step is the preprocessing before we can form our training set and start training the data. We will use `train_datagen` to incorporate images into our model during training.

```
1 # Create an ImageDataGenerator for test data (only rescaling)
2 test_datagen = ImageDataGenerator(rescale=1./255)
```

```
1 # Create training data flow
2 training_set = train_datagen.flow_from_directory(
3     image_dir, # Path to the main image directory
4     target_size=(64, 64), # Resize images to match input shape
5     batch_size=32,
6     class_mode='categorical', # Use 'categorical' for multi-class
7     classes=subfolders
8 )
```

Found 1578 images belonging to 3 classes.

```
1 # Create test data flow
2 test_set = test_datagen.flow_from_directory(
3     image_dir,
4     target_size=(64, 64),
5     batch_size=32,
6     class_mode='categorical', # Use 'categorical' for multi-class
7     classes=subfolders
8 )
```

Found 1578 images belonging to 3 classes.

For the test data, we do not want to augment it. So only rescaling is applied in the line

"test\_datagen = ImageDataGenerator(rescale=1./255) For the training set and test set, we use flow\_from\_directory to create data generators for both sets

The parameters are:

- image\_dir: directory containing the images
- target\_size: size for the resized images
- batch\_size: number of images to process in each batch
- class\_mode: classification problem type
- classes: classes from the list of subfolder names

We are now ready to train the CNN model. We use the training\_set that we just created. We set the other parameters: **steps\_per\_epoch**: number of batches to process in each epoch, the total number of training samples divided by the batch size **epochs**: number of times to iterate over the entire training dataset **validation\_data**: test\_set we just created used to validate the training **validation\_steps**: batches to process validation in each epoch.

```
1 # Fit the model - training
2 lr_scheduler = ReduceLROnPlateau(factor=0.5, patience=5, verbose=1)
3 cnn.fit(
4     training_set,
5     steps_per_epoch=len(training_set),
6     epochs=25,
7     validation_data=test_set,
8     validation_steps=len(test_set),
9     callbacks=[lr_scheduler],
10 )
```

Epoch 1/25

/usr/local/lib/python3.10/dist-packages/keras/src/trainers/data\_adapters/py\_data\_adapter.py:153: UserWarning: Your input ran out of data: <generator object PyDataAdapterGenerator at 0x7f8b1c1c1c1c>. You may need to use the argument `take\_steps` to specify the number of steps to run before stopping. (DataAdapters.py:153)

50/50 ————— 46s 845ms/step - accuracy: 0.5170 - f1\_metric: 0.31

Epoch 2/25

/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data: <generator object PyDataAdapterGenerator at 0x7f8b1c1c1c1c>. You may need to use the argument `take\_steps` to specify the number of steps to run before stopping. (DataAdapters.py:153)

50/50 ————— 4s 74ms/step - accuracy: 0.0000e+00 - f1\_metric: 0.00

Epoch 3/25

/usr/local/lib/python3.10/dist-packages/keras/src/callbacks/callback\_list.py:153: UserWarning: Your input ran out of data: <generator object PyDataAdapterGenerator at 0x7f8b1c1c1c1c>. You may need to use the argument `take\_steps` to specify the number of steps to run before stopping. (DataAdapters.py:153)

50/50 ————— 62s 509ms/step - accuracy: 0.6131 - f1\_metric: 0.51

Epoch 4/25

50/50 ————— 0s 684us/step - accuracy: 0.0000e+00 - f1\_metric: 0.00

Epoch 5/25

50/50 ————— 42s 519ms/step - accuracy: 0.6482 - f1\_metric: 0.61

Epoch 6/25

50/50 ————— 0s 635us/step - accuracy: 0.0000e+00 - f1\_metric: 0.00

Epoch 7/25

```

50/50 ————— 32s 61ms/step - accuracy: 0.6895 - f1_metric: 0.68
Epoch 8/25
50/50 ————— 1s 21ms/step - accuracy: 0.0000e+00 - f1_metric: 0.00
Epoch 9/25
50/50 ————— 37s 555ms/step - accuracy: 0.6748 - f1_metric: 0.60
Epoch 10/25
50/50 ————— 0s 670us/step - accuracy: 0.0000e+00 - f1_metric: 0.00
Epoch 11/25
50/50 ————— 28s 518ms/step - accuracy: 0.6724 - f1_metric: 0.60
Epoch 12/25
50/50 ————— 0s 652us/step - accuracy: 0.0000e+00 - f1_metric: 0.00
Epoch 13/25
50/50 ————— 40s 494ms/step - accuracy: 0.6975 - f1_metric: 0.69
Epoch 14/25
50/50 ————— 0s 690us/step - accuracy: 0.0000e+00 - f1_metric: 0.00
Epoch 15/25
50/50 ————— 41s 506ms/step - accuracy: 0.7079 - f1_metric: 0.68
Epoch 16/25
50/50 ————— 0s 9ms/step - accuracy: 0.0000e+00 - f1_metric: 0.00
Epoch 17/25
50/50 ————— 26s 496ms/step - accuracy: 0.7201 - f1_metric: 0.69
Epoch 18/25
50/50 ————— 0s 1ms/step - accuracy: 0.0000e+00 - f1_metric: 0.00
Epoch 19/25
50/50 ————— 26s 498ms/step - accuracy: 0.7111 - f1_metric: 0.69
Epoch 20/25
50/50 ————— 1s 20ms/step - accuracy: 0.0000e+00 - f1_metric: 0.00
Epoch 21/25
50/50 ————— 40s 507ms/step - accuracy: 0.6937 - f1_metric: 0.68
Epoch 22/25
50/50 ————— 0s 1ms/step - accuracy: 0.0000e+00 - f1_metric: 0.00
Epoch 23/25
50/50 ————— 41s 511ms/step - accuracy: 0.7343 - f1_metric: 0.73
Epoch 24/25
50/50 ————— 0s 683us/step - accuracy: 0.0000e+00 - f1_metric: 0.00
Epoch 25/25
50/50 ————— 26s 502ms/step - accuracy: 0.7105 - f1_metric: 0.71
<keras.src.callbacks.history.History at 0x7f5bd88577c0>

```

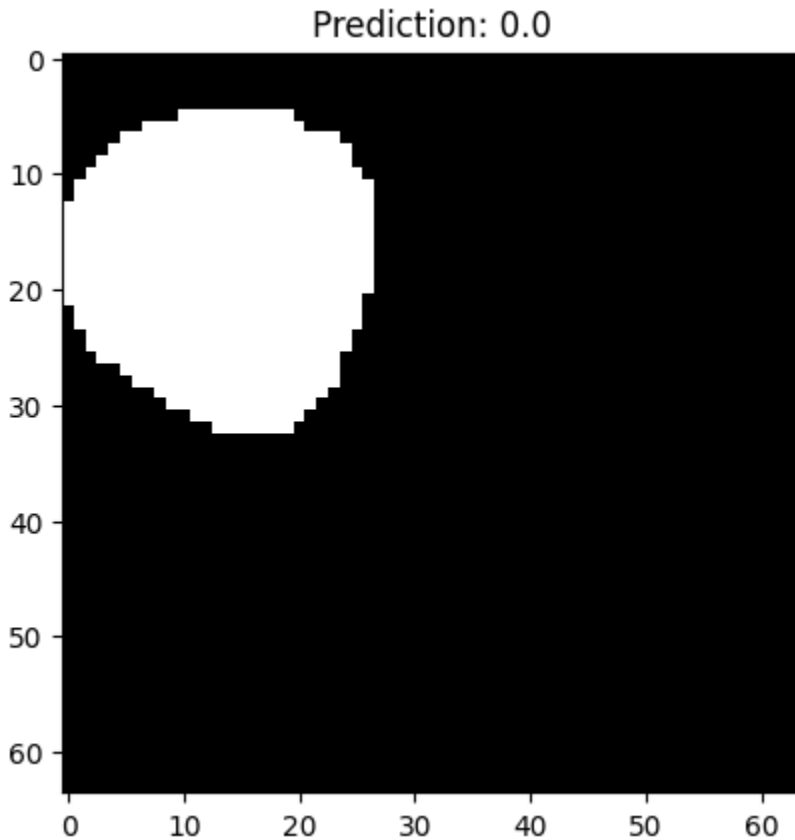
```

1 #make single prediction
2 from keras.preprocessing import image
3 import numpy as np
4
5 # Choose an image path
6 image_file_path = image_path[0]
7
8 test_image = image.load_img(image_file_path, target_size = (64, 64)) # Load us
9 test_image = image.img_to_array(test_image)
10
11
12 test_image = np.expand_dims(test_image, axis=0) # add a batch dimension
13 result = cnn.predict(test_image)
14

```

```
1 # Display the image using plt.imshow
2 plt.imshow(test_image[0]) # Display the first image in the batch
3 plt.title("Prediction: " + str(result[0][0])) # Add a title with the prediction
4 plt.show()
```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with



We select an image path for prediction with `image_file_path`. Then `test_image` loads the image, converts it to an array and adds a batch dimension. The result is predicted by the test image with the trained CNN model.

```
1 training_set.class_indices
{'benign': 0, 'malignant': 1, 'normal': 2}
```

```
1 if result[0][0] == 1:
2     prediction = 'benign',
3 elif result[0][0] == 2:
4     prediction = 'normal'
5 else:
6     prediction = 'malignant'
```

```
1 print(prediction)
```

```
    malignant
```

To access the dictionary that maps the class names (benign, normal, malignant) to the indices of the dictionary they are stored in, we use `training_set.class_indicies`.


We then have a chain of if, elif, else statements to interpret those indices and return the predicted class through the variable called `prediction`.


```
1 loss, accuracy, precision, recall, f1 = cnn.evaluate(test_set, steps=len(test_
2 print("Test Loss:", loss)
3 print("Test Accuracy:", accuracy)
4 print("Test Precision:", precision)
5 print("Test Recall:", recall)
6 print("Test F1-score:", f1)
```

```
50/50 ————— 10s 208ms/step - accuracy: 0.7071 - f1_metric: 0.70
Test Loss: 0.6331545114517212
Test Accuracy: 0.7116603255271912
Test Precision: 0.739130437374115
Test Recall: 0.7002534866333008
Test F1-score: 0.7189056873321533
```

```
1 from sklearn.metrics import confusion_matrix
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4 import numpy as np
5
6 y_pred = cnn.predict(test_set, steps=len(test_set))
7 y_pred_classes = np.argmax(y_pred, axis=1) # Convert pto class labels
8 cm = confusion_matrix(test_set.classes, y_pred_classes)
9
10 sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
11 plt.xlabel('Predicted')
12 plt.ylabel('Actual')
13 plt.title('Confusion Matrix')
14 plt.show()
```

```
50/50 ————— 10s 196ms/step
```





Analysis:

Accuracy: 74.4% The test accuracy is pretty good. The model is correctly classifying 74% of images in the test set.

Precision: 77.3% precise. Images predicted to belong to a specific class are about 77.3% correct.

Recall: 71.0% of the actual instances for each class are correctly identified

F1: 73.8% this number represents the difference between precision and recall.

Loss: 0.566% the loss function is actually quite low for multiple classification

The model's performance has a decent accuracy and demonstrates a good balance between precision (higher threshold) and recall(sensitivity, lower threshold). The appropriate loss function and output format is crucial and took some trial and error.

References

Dropout:

- <https://saturncloud.io/blog/where-to-add-dropout-in-neural-network/#2>

F1-score:

- [https://www.tensorflow.org/addons/api\\_docs/python/tfa/metrics/F1Score](https://www.tensorflow.org/addons/api_docs/python/tfa/metrics/F1Score)
- <https://stackoverflow.com/questions/68596302/f1-score-metric-per-class-in-tensorflow>

ImageDataGenerator:

- [https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image/](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/)

[https://www.tensorflow.org/api\\_docs/python/tf/keras/preprocessing/image/ImageDataGenerator](https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/image/ImageDataGenerator)

ReduceLROnPlateau:

- [https://www.tensorflow.org/api\\_docs/python/tf/keras/callbacks/ReduceLROnPlateau](https://www.tensorflow.org/api_docs/python/tf/keras/callbacks/ReduceLROnPlateau)

Loss Function, multiclass:

- <https://discuss.pytorch.org/t/loss-function-for-multi-class-with-probabilities-as-output/60866>
- [https://medium.com/@nghihuyinh\\_37300/understanding-loss-functions-for-classification-81c19ee72c2a](https://medium.com/@nghihuyinh_37300/understanding-loss-functions-for-classification-81c19ee72c2a)
- <https://stackoverflow.com/questions/59336899/which-loss-function-and-metrics-to-use-for-multi-label-classification-with-very>

Metrics:

- [https://scikit-learn.org/dev/modules/generated/sklearn.metrics.multilabel\\_confusion\\_matrix.html](https://scikit-learn.org/dev/modules/generated/sklearn.metrics.multilabel_confusion_matrix.html)
- <https://www.kaggle.com/code/kmkarakaya/multi-label-model-evaluation>
- <https://learn.microsoft.com/en-us/azure/machine-learning/how-to-understand-automated-ml?view=azureml-api-2>
- <https://developers.google.com/machine-learning/crash-course/classification/accuracy-precision-recall>
- <https://arize.com/blog-course/f1-score/>
- <https://stackoverflow.com/questions/55984768/i-am-having-trouble-calculating-the-accuracy-recall-precision-and-f1-score-for>
- <https://www.labelf.ai/blog/what-is-accuracy-precision-recall-and-f1-score>
- <https://medium.com/analytics-vidhya/precision-recall-tradeoff-79e892d43134>

```
1 # Modified code from lab
2 # Importing necessary libraries for data visualization, data handling, and mac
3 import matplotlib.pyplot as plt # For creating visualizations (plots)
4 import pandas as pd # For data manipulation and analysis (DataFrames)
5 import numpy as np # For numerical computations and working with arrays
6 import seaborn as sns # For statistical data visualization
7 import warnings # For handling warnings in the code
8 # Ensuring compatibility with Python 2.x for print function if needed
9 from __future__ import print_function # To use the print() function from Pyth
10
11 # Importing necessary modules from Keras for building neural networks
12 from keras.models import Model # For defining and training machine learning m
13 from keras.layers import Dense, Input # For defining layers in the neural net
14 from keras.regularizers import l1 # For L1 regularization to prevent overfitt
15 from keras.optimizers import Adam # Adam optimizer for model training
```



```
15 from keras.optimizers import Adam # Adam optimizer for model training
```

```
1 # Changed from Lab session
2 # Define the dimensions of the input, hidden layers, and encoded representatio
3 input_size = 64 * 64 * 3 # Input size (for 64x64x3 images flattened into a 40
4 hidden_size = 128 # Number of neurons in the hidden layers
5 code_size = 32 # Size of the encoded representation (bottleneck)
6
7 # Define the input layer with the specified input size
8 input_img = Input(shape=(input_size,))
9
10 # Define the first hidden layer with 128 neurons and ReLU activation
11 hidden_1 = Dense(hidden_size, activation='relu')(input_img)
12
13 # Define the encoded layer (code) with 32 neurons and ReLU activation
14 code = Dense(code_size, activation='relu')(hidden_1)
15
16 # Define the second hidden layer, which reconstructs the input size
17 hidden_2 = Dense(hidden_size, activation='relu')(code)
18
19 # Define the output layer to reconstruct the input image, using softmax activa
20 output_img = Dense(input_size, activation='softmax')(hidden_2)
21
22 # Create the autoencoder model, specifying input and output layers
23 autoencoder = Model(input_img, output_img)
24
25 # Compile the autoencoder model using the Adam optimizer and categorical cross
26 autoencoder.compile(optimizer='adam', loss='categorical_crossentropy')
```

```
1 # Modified code from lab
2 # X_train Load and preprocess images before training
3 X_train_images = []
4 for img_path in X_train:
5     # Find the subfolder where the image belongs
6     for subfolder in subfolders:
7         if img_path in os.listdir(os.path.join(image_dir, subfolder)):
8             full_path = os.path.join(image_dir, subfolder, img_path)
9             break
10
11 img = keras.utils.load_img(full_path, target_size=(64, 64)) # Load image
12 img = keras.utils.img_to_array(img) # Convert to array
13 img = img / 255.0 # Rescale pixel values
14 img = img.flatten() # Flatten the image before appending
15 X_train_images.append(img)
16 X_train_images = np.array(X_train_images)
17
18
```

```
1 # Modified code from lab
```

```

2 # Train the autoencoder on the training data, using the same data for both inp
3 autoencoder.fit(X_train_images, X_train_images, epochs=10) # Training for 3 e

```

```

Epoch 1/10
40/40 ————— 4s 63ms/step - loss: 21050.9980
Epoch 2/10
40/40 ————— 2s 45ms/step - loss: 243634.1094
Epoch 3/10
40/40 ————— 2s 42ms/step - loss: 1609054.2500
Epoch 4/10
40/40 ————— 2s 40ms/step - loss: 4929993.0000
Epoch 5/10
40/40 ————— 3s 41ms/step - loss: 11525234.0000
Epoch 6/10
40/40 ————— 2s 42ms/step - loss: 22911488.0000
Epoch 7/10
40/40 ————— 2s 61ms/step - loss: 38971996.0000
Epoch 8/10
40/40 ————— 3s 64ms/step - loss: 60547300.0000
Epoch 9/10
40/40 ————— 4s 42ms/step - loss: 89374304.0000
Epoch 10/10
40/40 ————— 3s 41ms/step - loss: 126374360.0000
<keras.src.callbacks.history.History at 0x7f5bcb1eb010>

```

```

1 # Modified code from lab
2 # X_test Load and preprocess images before prediction
3 X_test_images = []
4 for img_path in X_test:
5     # Find the subfolder where the image belongs
6     for subfolder in subfolders:
7         if img_path in os.listdir(os.path.join(image_dir, subfolder)):
8             full_path = os.path.join(image_dir, subfolder, img_path)
9             break
10
11 img = keras.utils.load_img(full_path, target_size=(64, 64)) # Load image
12 img = keras.utils.img_to_array(img) # Convert to array
13 img = img / 255.0 # Rescale pixel values
14 img = img.flatten() # Flatten the image before appending
15 X_test_images.append(img)
16 X_test_images = np.array(X_test_images)
17

```

```

1 # Modified code from lab
2 # Generate the reconstructed (decoded) images from the autoencoder
3 decoded_imgs = autoencoder.predict(X_test_images)

```

```

10/10 ————— 0s 13ms/step

```

```

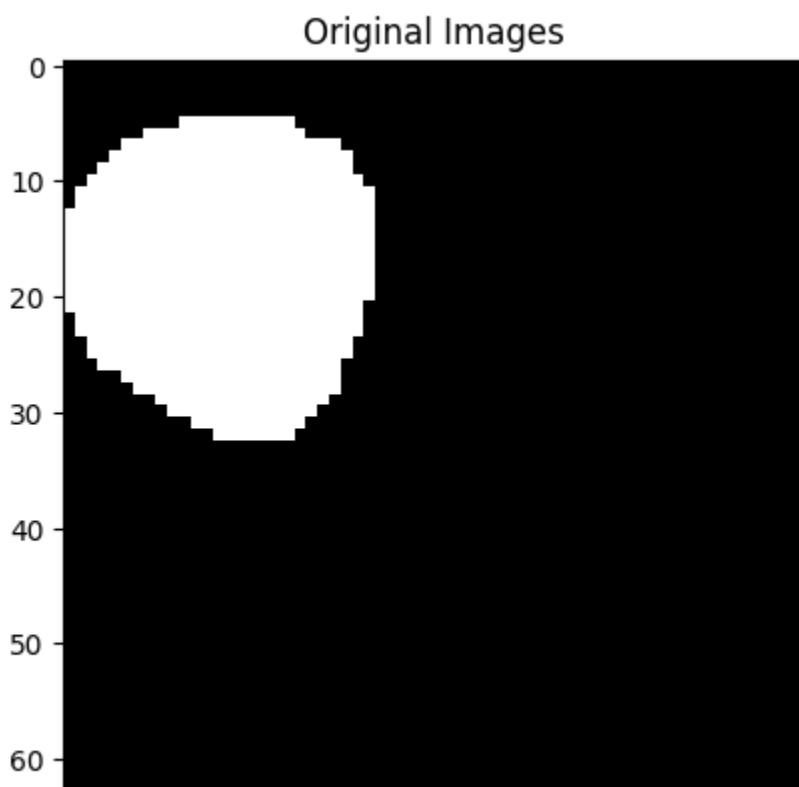
1 # Modified code from lab

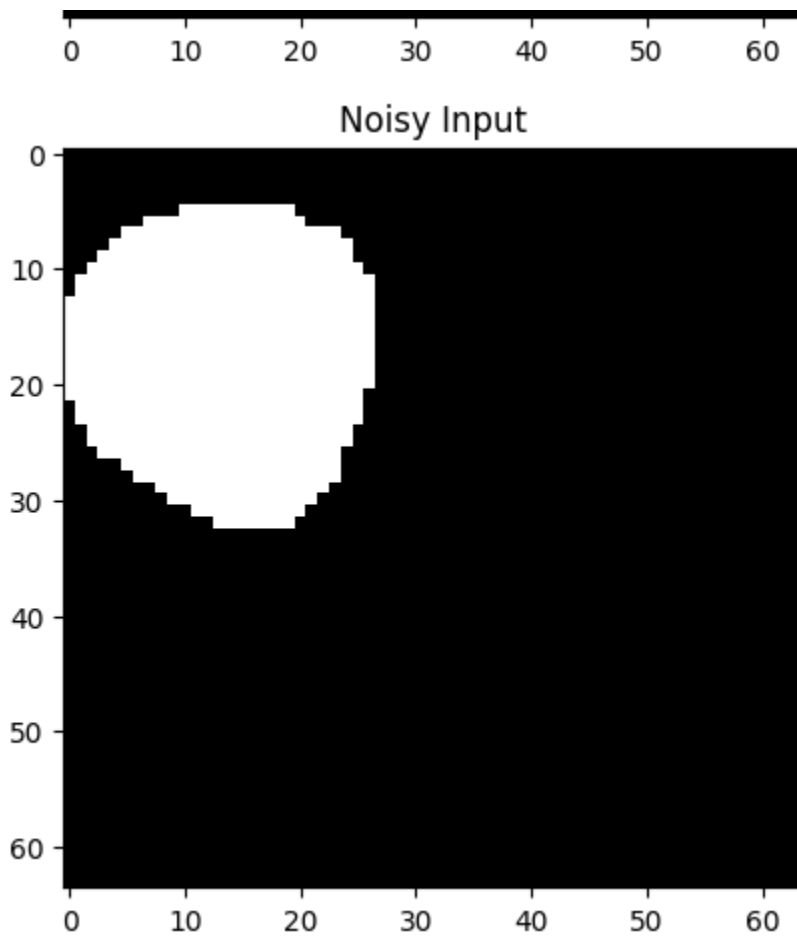
```

```

2 # Define the noise factor (standard deviation of the Gaussian noise to add to
3 noise_factor = 0.4
4
5 # X_train Load and preprocess images before training
6 X_train_float = []
7 for img_path in X_train:
8     X_train_float.append(img) # Append the processed image data
9 X_train_float = np.array(X_train_float) # Convert to NumPy array
10
11 # X_test Load and preprocess images before training
12 X_test_float = []
13 for img_path in X_test:
14     X_test_float.append(img) # Append the processed image data
15 X_test_float = np.array(X_test_float) # Convert to NumPy array
16
17
18 # Add random Gaussian noise to the training and test images
19 X_train_noisy = X_train_float + noise_factor * np.random.normal(size=X_train_f
20 X_test_noisy = X_test_float + noise_factor * np.random.normal(size=X_test_floa
21
22 # Clip the noisy images to ensure pixel values stay within the range [0, 1]
23 x_train_noisy = np.clip(X_train_noisy, 0.0, 1.0)
24 x_test_noisy = np.clip(X_test_noisy, 0.0, 1.0)
25 plt.imshow(X_test_float[0].reshape(64, 64, 3)) # Reshape to (64, 64, 3) for a
26 plt.title('Original Images')
27 plt.show()
28 plt.imshow(X_test_float[0].reshape(64, 64, 3)) # Reshape to (64, 64, 3) for a
29 plt.title('Noisy Input')
30 plt.show()

```






```
1 # Modified code from lab
2 # Set the number of example digits to display
3 n = 5
4
5 # Create a figure to display the original and noisy images
6 plt.figure(figsize=(10, 4.5))
7
8 # Loop over the first n images to plot their original and noisy versions
9 for i in range(n):
10     # Load and preprocess the original image
11     img_path = X_test[i]
12     for subfolder in subfolders:
13         if img_path in os.listdir(os.path.join(image_dir, subfolder)):
14             full_path = os.path.join(image_dir, subfolder, img_path)
15             break
16     original_img = keras.utils.load_img(full_path, target_size=(64, 64))
17     original_img = keras.utils.img_to_array(original_img) / 255.0 # Rescale
18
19     # Plot the original image in the top row
20     ax = plt.subplot(2, n, i + 1)
21     plt.imshow(original_img) # Display the loaded and preprocessed image
22     plt.gray()
23     ax.get_xaxis().set_visible(False)
```

```

24 ax.get_yaxis().set_visible(False)
25 if i == n // 2:
26     ax.set_title('Original Images')
27
28 # Plot the noisy image in the bottom row
29 ax = plt.subplot(2, n, i + 1 + n)
30 plt.imshow(X_test_noisy[i].reshape(64, 64, 3))
31 plt.gray()
32 ax.get_xaxis().set_visible(False)
33 ax.get_yaxis().set_visible(False)
34 if i == n // 2:
35     ax.set_title('Noisy Input')
36
37 # Display the plot
38 plt.show()

```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow wi  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow wi  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow wi  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow wi  
 WARNING:matplotlib.image:Clipping input data to the valid range for imshow wi



```

1 # Modified code from lab
2 # Define the input size, hidden layer size, and code (bottleneck) size for the
3 input_size = 64 * 64 * 3
4 hidden_size = 128 # Number of neurons in the hidden layers
5 code_size = 32 # Size of the encoded representation (bottleneck layer)
6
7 # Define the input layer with the specified input size
8 input_img = Input(shape=(input_size,))

```

```

9
10 # First hidden layer with ReLU activation
11 hidden_1 = Dense(hidden_size, activation='relu')(input_img)
12
13 # Code (bottleneck) layer with 32 neurons and ReLU activation
14 code = Dense(code_size, activation='relu')(hidden_1)
15
16 # Second hidden layer with ReLU activation, taking input from the code layer
17 hidden_2 = Dense(hidden_size, activation='relu')(code)
18
19 # Output layer, reconstructing the original input, with a softmax activation
20 output_img = Dense(input_size, activation='softmax')(hidden_2)
21
22 # Create the autoencoder model with the input and output layers
23 autoencoder_noisy = Model(input_img, output_img)
24
25 # Compile the model using the Adam optimizer and categorical cross-entropy loss
26 autoencoder_noisy.compile(optimizer='adam', loss='categorical_crossentropy')
27
28 X_train_images = []
29 for img_path in X_train:
30     # Find the subfolder where the image belongs
31     for subfolder in subfolders:
32         if img_path in os.listdir(os.path.join(image_dir, subfolder)):
33             full_path = os.path.join(image_dir, subfolder, img_path)
34             break # Exit the inner loop once the image is found
35
36     # Load the image, convert to array, and rescale
37     img = keras.utils.load_img(full_path, target_size=(64, 64))
38     img = keras.utils.img_to_array(img)
39     img = img / 255.0
40     img = img.flatten() # If your autoencoder requires flattened input
41     X_train_images.append(img) # Append the processed image data
42
43 # Convert the list of images into a NumPy array
44 X_train_images = np.array(X_train_images)
45
46 # Add random Gaussian noise to the training images
47 X_train_noisy = X_train_images + noise_factor * np.random.normal(size=X_train_
48
49 # Clip the noisy images to ensure pixel values stay within the range [0, 1]
50 X_train_noisy = np.clip(X_train_noisy, 0.0, 1.0)
51
52 # Train the autoencoder model
53 # Use X_train_images (original images) as the target for reconstruction
54 autoencoder_noisy.fit(X_train_noisy, X_train_images, epochs=10) # Changed X_tr
55

```

Epoch 1/10

**40/40**  3s 41ms/step - loss: 33279.4453

Epoch 2/10

```

Epoch 2/10
40/40 ————— 2s 41ms/step - loss: 547962.6250
Epoch 3/10
40/40 ————— 3s 42ms/step - loss: 3152904.2500
Epoch 4/10
40/40 ————— 2s 41ms/step - loss: 10835492.0000
Epoch 5/10
40/40 ————— 2s 51ms/step - loss: 26716880.0000
Epoch 6/10
40/40 ————— 3s 63ms/step - loss: 53235696.0000
Epoch 7/10
40/40 ————— 2s 44ms/step - loss: 92860528.0000
Epoch 8/10
40/40 ————— 2s 42ms/step - loss: 146243328.0000
Epoch 9/10
40/40 ————— 2s 41ms/step - loss: 212786784.0000
Epoch 10/10
40/40 ————— 2s 42ms/step - loss: 299673248.0000
<keras.src.callbacks.history.History at 0x7f5bcabd8a90>

```

```

1 # Modified code from lab
2 # Set the number of images to display
3 n = 5
4
5 # Create a figure to display the original, noisy, and autoencoder output image
6 plt.figure(figsize=(10, 7))
7
8 # Generate the denoised (reconstructed) images using the trained autoencoder
9 # Changed autoencoder to autoencoder_noisy
10 images = autoencoder_noisy.predict(X_test_noisy)
11
12 # Loop over the first n images to plot their original, noisy, and reconstructe
13 for i in range(n):
14     # Load and preprocess the original image
15     img_path = X_test[i]
16     for subfolder in subfolders:
17         if img_path in os.listdir(os.path.join(image_dir, subfolder)):
18             full_path = os.path.join(image_dir, subfolder, img_path)
19             break
20     # Load the original image and reshape for display
21     original_img = keras.utils.load_img(full_path, target_size=(64, 64))
22     original_img = keras.utils.img_to_array(original_img) / 255.0 # Rescale
23
24     # Plot the original image in the top row
25     ax = plt.subplot(3, n, i + 1) # Create a subplot for the original image
26     plt.imshow(original_img) # Display the loaded and preprocessed image
27     plt.gray() # Use grayscale for the image
28     ax.get_xaxis().set_visible(False) # Hide the x-axis for a cleaner plot
29     ax.get_yaxis().set_visible(False) # Hide the y-axis for a cleaner plot
30     if i == n // 2: # Add a title in the middle
31         ax.set_title('Original Images')
32

```

```

33 # Plot the noisy image in the middle row
34 ax = plt.subplot(3, n, i + 1 + n) # Create a subplot for the noisy image
35 plt.imshow(x_test_noisy[i].reshape(64, 64, 3)) # Reshape and display the
36 plt.gray() # Use grayscale for the image
37 ax.get_xaxis().set_visible(False) # Hide the x-axis
38 ax.get_yaxis().set_visible(False) # Hide the y-axis
39 if i == n // 2: # Add a title in the middle
40     ax.set_title('Noisy Input')
41
42 # Plot the autoencoder's output (denoised) image in the bottom row
43 ax = plt.subplot(3, n, i + 1 + 2*n) # Create a subplot for the reconstruc
44 plt.imshow(images[i].reshape(64, 64, 3)) # Reshape and display the recons
45 plt.gray() # Use grayscale for the image
46 ax.get_xaxis().set_visible(False) # Hide the x-axis
47 ax.get_yaxis().set_visible(False) # Hide the y-axis
48 if i == n // 2: # Add a title in the middle
49     ax.set_title('Autoencoder Output')
50
51 # Display the plot
52 plt.show()

```

10/10 ————— 0s 14ms/step


Original Images



Noisy Input



Autoencoder Output





Autoencoder class is made by subclassing `tf.keras.Model`. This auto encoder is used to denoise the images from the breast ultrasound image dataset.

### 1. PreProcessing:

- load the image from the previously created (`image_dir`)
- resize the image, may not be necessary, just in case
- normalize the image, this was done previously in preprocessing for CNN, we are doing it again here
- flatten to 1D array

### 2. Autoencoder Architecture

- Keras function API
- layers: input layer, two hidden layers (encoder, decoder), and an output layer
- hidden layers use ReLU
- output layer uses softmax

### 3. Training

- autoencoder trained on preprocessed images (`X_train_images`)
- trained for 10 epochs with Adam as the optimizer and categorical cross-entropy as the loss function

### 4. Denoising

- using the code from the autoencoder lab, the Gaussian noise is added to the original images
- the autoencoder is trained and then denoises the images

### 5. Visualization

- the denoised images are plotted. The denoised images do not look like the original. Perhaps this has to do with the fact that the printed images are from the masked dataset. With more time and experience perhaps the programmer could do a better job with the dataset.

### References AutoEncoder:

- <https://www.pyquantnews.com/the-pyquant-newsletter/use-autoencoders-to-create-feature-embeddings>
- <https://www.tensorflow.org/tutorials/generative/autoencoder>

- 
- <https://pierre-schwartz.medium.com/introduction-aux-auto-encodeurs-61e8d74660f3>
  - <https://hex.tech/blog/autoencoders-for-feature-selection/>
  - <https://www.mathworks.com/matlabcentral/answers/1619715-how-do-i-find-the-features-extracted-from-input-data-set-using-auto-encoder>
  - <https://stackoverflow.com/questions/55920986/how-to-feed-time-series-data-into-an-autoencoder-network-for-feature-extraction>
  - <https://blog.keras.io/building-autoencoders-in-keras.html>
  - [https://mallahyari.github.io/ml\\_tutorial/autoencoder/](https://mallahyari.github.io/ml_tutorial/autoencoder/)